

Team Checkpoint 2

Forest Cover Team B

Annie Condon, Matt Hayden, Matt Robertson, Yvette Gonzalez

Contents

1	Introduction	3
2	The Modeling Problem	3
2.1	Evaluating Classification Models	3
3	The Data	4
3.1	Response Variable	4
3.2	Continuous Variables	5
3.3	Binary Variables	5
4	Exploratory Data Analysis	7
5	Data Preparation	15
5.1	Transformations	15
5.2	Splitting the Data	17
6	Exploratory Models - Feature Selection	17
6.1	Linear Discriminant Analysis	17
6.2	Decision Tree	17
7	Conclusions from EDA and Data Preparation	18

1 Introduction

The U.S. Forest Service relies on an accurate understanding of its forests composition in order to best protect and manage the forest land. Conducting accurate inventory of forest composition by direct observation or remotely sensed data is often too expensive and time consuming to do at large-scale. Predictive analytics can be employed to use the results of a small-scale survey to create a model that can be applied across a large region, using descriptive features extracted from maps of the area.

In this paper, our objective is to predict the forest cover type given a set of cartographic features and a variety of multiclass classification models. Our models will be evaluated using predictive accuracy.

2 The Modeling Problem

Our modeling problem is to predict the forest cover type as a multiclass classification problem based on the associated features. A multiclass classification problem classifies instances into one of the more than two classes. Our forest cover type is defined as one of seven, mutually exclusive, forest cover type classes, shown in table 1 below.

Table 1: Cover Type Classes

Cover Type	Number of Observations
Spruce/Fir	211840
Lodgepole Pine	283301
Ponderosa Pine	35754
Cottonwood/Willow	2747
Aspen	9493
Douglas/Fir	17367
Krummholz	20510

Most classification algorithms can be applied, either directly or through slight modifications, to multiclass classification problems. Two different approaches exist for multi-class classification: algorithms that naturally permit the use of more than two classes and those that first reduce a multi-class problem to a collection of binary-class problems and then combine their predictions in various ways. The following is a list of algorithms we will consider for our problem:

- Discriminant Analysis
- Logistic Regression
- K-Nearest Neighbors
- Random Forests
- Neural Network
- Support Vector Machine

We will use ‘Kappa’ as a metric to optimize for tuning our parameters because this metric can improve the quality of the model for problems where there are a low percentage of samples in one class.

2.1 Evaluating Classification Models

We will evaluate a number of different models by applying them to a holdout ‘test’ dataset and assessing their performance. Because there is an imbalance in the classes to be predicted, we will consider alternatives to ‘accuracy’ to evaluate our final models. Using accuracy only to evaluate models can perform poorly in predicting the less frequent classes. We will instead look at two metrics that handle class imbalance: balanced accuracy and f1 score. Balanced accuracy is (sensitivity+specificity)/2, or the average accuracy over all

classes. F1 Score is $\text{precision} \times \text{recall} / (\text{precision} + \text{recall})$, the weighted average of precision and recall. The precision measures the accuracy of a predicted positive outcome and recall (sensitivity) measures the strength of the model to predict a positive outcome. Therefore, this score takes both false positives and false negatives into account.

3 The Data

Our data set consists of 581,012 observations of the 30 x 30 meter cells of forest and 54 features associated with each cell. The features are derived from 12 attributes, with area and soil type binarized so that there are 4 binary area designators and 40 binary soil type designators. There is no missing data. The feature descriptions are listed in the table below:

Table 2: Features

Feature	Descriptions
Elevation	Elevation in meters
Aspect	Aspect in degrees azimuth
Slope	Slope in degrees
HDist.Hydrology	Horizontal distance to nearest surface water feature in meters
VDist.Hydrology	Vertical distance to nearest surface water feature in meters
HDist.Roadway	Horizontal distance to nearest roadway in meters
Hillshade.9am	Hillshade index at 9am, summer solstice
Hillshade.12pm	Hillshade index at noon, summer solstice
Hillshade.3pm	Hillshade index at 3pm, summer solstice
HHDist.FirePoint	Horizontal Distance to nearest wildfire ignition points
Area	Wilderness area designation - 4 binary areas
SoilType	Soil Type designation - 40 binary values

3.1 Response Variable

Figure one shows the frequency of each class of cover type in our data set. 85 percent of the observations fall into classes 1 and 2, Spruce/Fir and Lodgepole Pine. Class 4, Cottonwood/Willow has the least amount of observations at 2,747.

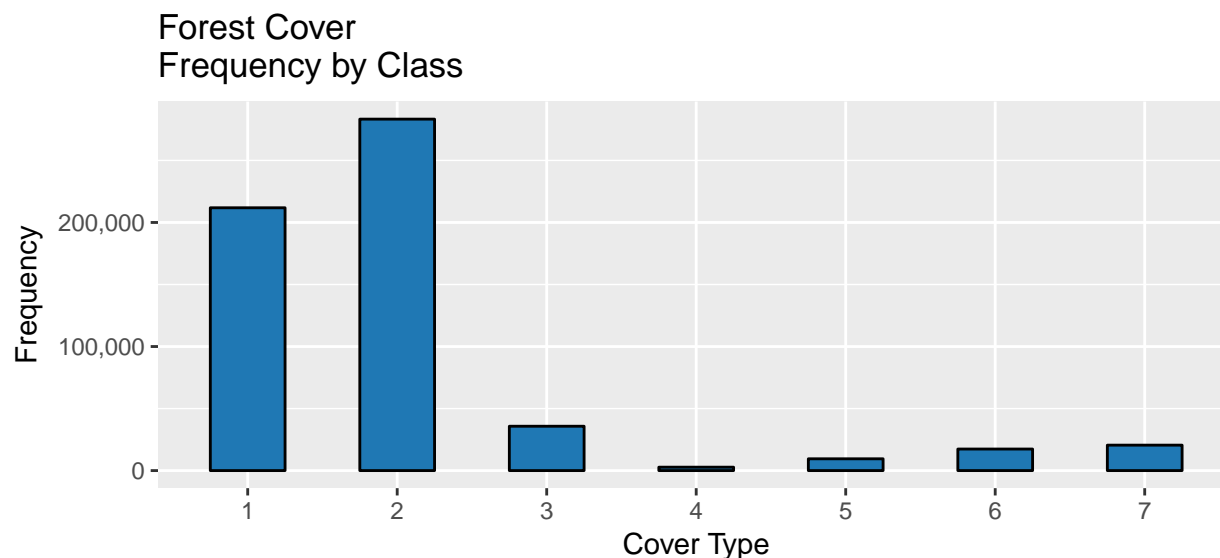


Figure 1: Frequency of each Cover Type Class

3.2 Continuous Variables

Table 3 includes some standard statistical measures of central tendency and variation for our continuous variables, for investigation of unusual values. We would first like to note that an Aspect value of 0 and 360 would both be equal to true north, so we should standardize that value. Next, we will point out that while a negative distance value for VDist.Hydrology may appear to be an error, it is in fact reasonable due to it being a vertical measurement from differing altitudes, making a negative distance possible. Our Hillshade values fall into the hillshade index integer value range of 0 to 255. Also, our horizontal distances all have minimum values of 0.

Table 3: Summary Statistics for Continuous Data

	mean	sd	median	min	max	nmiss	type
Elevation	2959.3653	279.9847	2996	1859	3858	0	Continuous
Aspect	155.6568	111.9137	127	0	360	0	Continuous
Slope	14.1037	7.4882	13	0	66	0	Continuous
HDist.Hydrology	269.4282	212.5494	218	0	1397	0	Continuous
VDist.Hydrology	46.4189	58.2952	30	-173	601	0	Continuous
HDist.Roadway	2350.1466	1559.2549	1997	0	7117	0	Continuous
Hillshade.9am	212.146	26.7699	218	0	254	0	Continuous
Hillshade.12pm	223.3187	19.7687	226	0	254	0	Continuous
Hillshade.3pm	142.5283	38.2745	143	0	254	0	Continuous
HDist.FirePoint	1980.2912	1324.1952	1710	0	7173	0	Continuous

3.3 Binary Variables

Our data's binary variables consist of 4 different wilderness area designators and 40 different soil type designators. Table 4 and 5 below represent the frequency of each of the areas and soil types, in descending order.

Table 4: Area Type Counts

Name	Count
Area1	5260796
Area3	253364
Area4	36968
Area2	29884

Name	Count
SoilType13	17431
SoilType38	15573
SoilType39	13806
SoilType11	12410
SoilType4	12396
SoilType20	9259
SoilType40	8750
SoilType2	7525
SoilType6	6575
SoilType3	4823
SoilType19	4021
SoilType17	3422
SoilType1	3031
SoilType16	2845
SoilType26	2589
SoilType18	1899
SoilType35	1891
SoilType34	1611
SoilType5	1597
SoilType9	1147
SoilType27	1086
SoilType28	946
SoilType21	838
SoilType14	599
SoilType25	474
SoilType37	298
SoilType8	179
SoilType36	119
SoilType7	105
SoilType15	3

4 Exploratory Data Analysis

Our next step is to explore the relationships in our data. We will begin by looking at boxplots of our scaled continuous variables in order to understand their relative distribution. We note that `Hillshade.12pm` and `VDist.Hydrology` have more pronounced skews than the other variables. We will need to investigate transformations if we use modeling techniques that can be negatively effected by outliers.




Figure 2: Boxplots of Scaled Continuous Variables

Our data exploration is informed by our statistical problem, which is one of multiclass classification. We will therefore be interested in exploring our predictors by each of our classes. The density plots below superimpose the density estimates for each variable by class, 1-7. We note that our variable Aspect shows multimodal distributions, indicating that it may have more than one grouping included. This is due to the value of 0 and 360 both being equal to true north, which we will have to address in data preparation. We will also note that several of our variables show distinct distributions by class, indicating that they would serve as a good predictor.

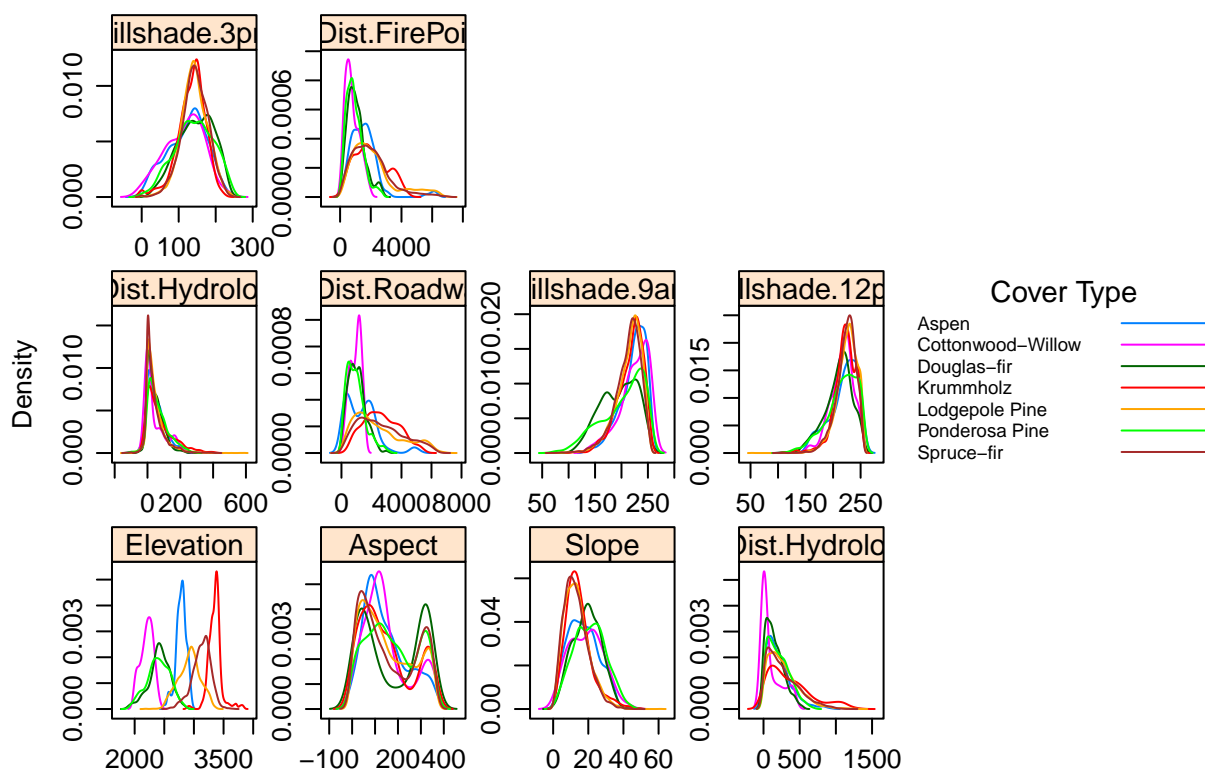


Figure 3: Density Plots of Continuous Variables by Class

Figure 4 shows density plots by each of the four wilderness areas. Several variable values appear to be distinguishable by wilderness area.

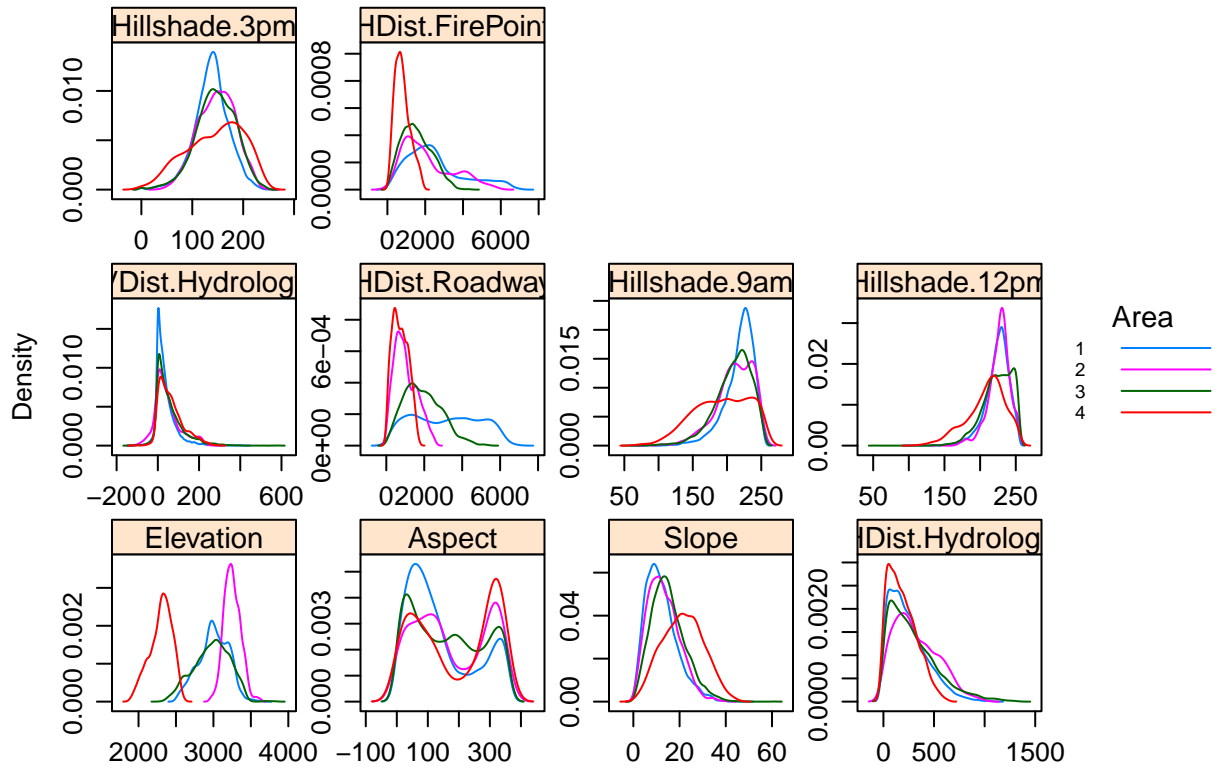


Figure 4: Density Plots of Continuous Variables by Area

Next, we will look at the correlations between the our predictor variables. Highly correlated predictors can have a negative effect on some of our modeling algorithms. Below is a correlation matrix with dark blue colors indicating a strong positive correlation and dark red colors indicating a strong negative correlation. We see the following variables with higher correlations that should be investigated further:

High Positive Collinearity

- VDist.Hydrology and HDist.Hydrology
- Aspect and Hillshade.3pm
- Hillshade.3pm and Hillshade.12pm

High Negative Collinearity

- Aspect and Hillshade.9am
- Slope and Hillshade.12pm
- Hillshade.3pm and Hillshade.9am




Figure 5: Correlation Matrix

Figure 6 shows a barchart of our 40 different soil types which shows the proportion of frequency of each class. We can see that soil types 1-7 have similar proportions, as do 19-33 and 38-40. A few soil types have only one cover type class, making them especially good predictors.




Figure 7 shows a barchart of our 4 different wilderness areas with the proportion of frequency of each class. We can see that the class composition in area 4 is distinguishable from the other areas, making it a good predictor.




Figure 7: Area Barchart

Next, we look at the horizontal distance to the roadway, by class. We can see that most observations in classes 5-7, Aspen, Cottonwood-Willow and Douglas-fir are closest to the roadways.




Figure 8: Distance to Roadway, by class

Figure 9 shows horizontal distance to the water, by class. Again, we see that most observations in classes 5-7 are closest to a water source.




Figure 9: Distance to water, by class

Figure 10 shows the contrasting relationship between Hillshade.9am and Hillshade.3pm.




Figure 10: Hillshade.9am and Hillshade.3pm


5 Data Preparation

5.1 Transformations


The distribution for Hillshade.12pm was noticeably skewed in the boxplots above, so we create a log transformation of Hillshade.12pm and for later use in our models. Additionally we created an interaction variable using Vist.Hydrology and HDistHydrology. A linear distance variable was also created using the vertical and horizontal distance to hydrology variables. Additionally, soil types were grouped into Climatic and Geologic zones as a potential means to reduce the amount of SoilType variables used for various modeling techniques.

Furthermore, we attempt to reduce the variable complexity of 40 soil variables with principal component analysis. We found that 90.7% of the variance can be explained with 16 variables instead of 40. Similar analysis using 6 of 10 Soil Zones explains 99.2% of variation. This may help with modeling efforts the rely on reducing model complexity.

Soil Type PCA Importance Plot



Soil Zone PCA Importance Plot



The correlation plot below also shows high correlations between Hillshade, Aspect, and Slope. By performing PCA analysis we are able to reduce 5 variables down to 3, with 0 correlation and explaining 99.7% of the variance. This is highly likely to be useful for models that need to minimize highly correlated variables.




Figure 11: PCA Correlation

Several of the multi-classification models we will rely on accounting for variables with near zero variance. Neural networks, logistic regression and KNN models prefer that no near zero variance variables are in the dataset. In this dataset, 32 of the 40 soil variables have near zero variance. We will look to principal component analysis, combining predictors, and removing predictors to remedy this issue.

5.2 Splitting the Data

We have split our data into 70/30 training-testing set so that we could evaluate the performance of our models.

6 Exploratory Models - Feature Selection

Next we will run a few exploratory models to help us understand the relationships in our data and help to evaluate the most important variables for feature selection.

6.1 Linear Discriminant Analysis

Linear Discriminant Analysis tries to find a linear combination of the predictors that gives maximum separation between the centers of the data while at the same time minimizing the variation within each group of data. We ran an exploratory model on the data that returned the linear combinations of the original variables that were created to distinguish between the classes. The variables with large absolute values in the scalings are more likely to be influential. For this data, Elevation, Area4 and trans.zone27 seem to be among the important variables. The first discriminant function, LD1, achieves 73.85% of the separation of classes, with the second discriminant function, LD2, improving the separation by 18.19%. Therefore, to achieve a good separation of the classes, we should use both of the first two discriminant functions.

6.2 Decision Tree

The tree plot below shows the result of fitting a decision tree algorithm to all of our data, the purpose of which is to draw insights regarding predictor variables that could be most effective in building a predictive model. Each node in the tree shows the predicted class, the predicted probability of each class and the percentage of observations in the node. Our tree uses the variables Elevation, Trans.zone87 and Area3 to predict our classes.

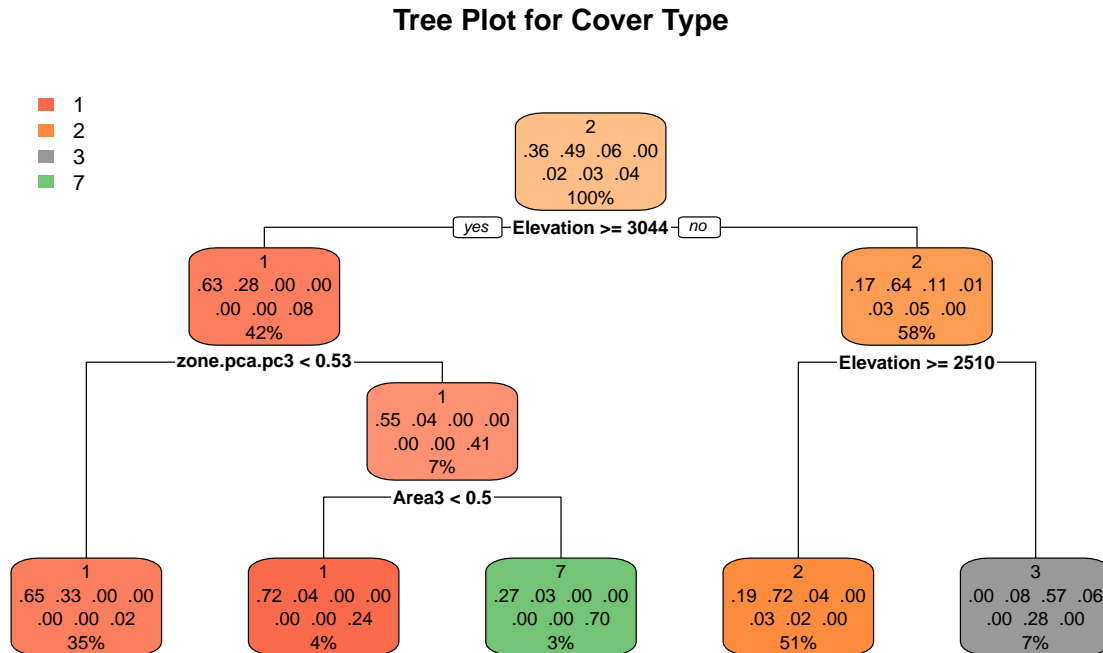


Figure 12: Tree Plot

7 Conclusions from EDA and Data Preparation

Our initial analysis for our forest cover type prediction problem included defining our modeling problem, doing a data quality and inventory check and performing a preliminary exploratory data analysis. The results of our data quality check showed that we had no missing data and gave us a cursory understanding of our data set.

Our preliminary exploratory analysis revealed some potentially useful predictors and helped us to understand the relationships in our data. Principal components analysis helped us narrow our soil variables down from 40 to 16 variables. Additionally, we prepared the data by creating a subset of 70% of the data to train our models on, and a validation set of 30% of the data to test those models on. Our next step will be to build our selected models using various subsets of our data, guided by our exploratory data analysis.